

Policies

- Due 9 PM, March 2nd, via Gradescope.
- You are free to collaborate on all of the problems, subject to the collaboration policy stated in the syllabus.
- In this course, we will be using Google Colab for code submissions. You will need a Google account.
- This set uses PyTorch, a Python package for neural networks. We recommend using Google Colab, which comes with PyTorch already installed.

Submission Instructions

- Submit your report as a single .pdf file to Gradescope, under "Set 6 Report".
- In the report, include any images generated by your code along with your answers to the questions.
- Submit your code by sharing a link in your report to your Google Colab notebook for each problem (see naming instructions below). Make sure to set sharing permissions to at least "Anyone with the link can view". Links that can not be run by TAs will not be counted as turned in. Check your links in an incognito window before submitting to be sure.
- For instructions specifically pertaining to the Gradescope submission process, see https://www.gradescope.com/get_started#student-submission.

Google Colab Instructions

For each notebook, you need to save a copy to your drive.

1. Open the github preview of the notebook, and click the icon to open the colab preview.
2. On the colab preview, go to File → Save a copy in Drive.
3. Edit your file name to "lastname_firstname_set_problem", e.g. "yue_yisong_set6_prob2.ipynb"

1 Class-Conditional Densities for Binary Data [25 Points, 8 EC Points]

This problem will test your understanding of probabilistic models, especially Naive Bayes. Consider a generative classifier for C classes, with class conditional density $p(x|y)$ and a uniform class prior $p(y)$. Suppose all the D features are binary, $x_j \in \{0,1\}$. If we assume all of the features are conditionally independent, as in Naive Bayes, we can write:

$$p(x | y = c) = \prod_{j=1}^D p(x_j | y = c)$$

This requires storing DC parameters.

Now consider a different model, which we will call the ‘full’ model, in which all the features are fully dependent.

Problem A [9 points]: Use the chain rule of probability to factorize $p(x | y)$, and let $\theta_{xjc} = p(x_j | x_1, \dots, x_{j-1}, y = c)$. Assuming we store each θ_{xjc} , how many parameters are needed to represent this factorization? Use big-O notation.

Solution A: Using chain rule, we can write this probability as:

$$p(x | y = c) = p(x_D | x_{D-1}, \dots, x_1, y = c) p(x_{D-1} | x_{D-2}, \dots, x_1, y = c) \dots p(x_1 | y = c) \quad (1)$$

In terms of the variable θ_{xjc} :

$$p(x | y = c) = \theta_{x_D c} \theta_{x_{D-1} c} \dots \theta_{x_1 c} \quad (2)$$

or:

$$p(x | y = c) = \prod_j \theta_{x_j c} \quad (3)$$

Thus, we need $\mathcal{O}(D)$ parameters for this factorisation.

Given each θ_{xjc} must store the probability for every prefix of length $j - 1$ (i.e. 2^{j-1} parameters), the total will be $\mathcal{O}\left(C \sum_{j=0}^{D-1} 2^j\right) = \mathcal{O}(C2^D)$. Thus, we need to store $\mathcal{O}(C2^D)$ number of parameters.

Problem B [8 points]: Assume we did no such factorization, and just used the joint probability $p(x | y = c)$. How many parameters would we need to estimate in order be able to compute $p(x|y = c)$ for arbitrary x and c ? How does this compare to your answer from the previous part? Again, use big-O notation.

Solution B: If we did not do this factorisation, then we would need to store a parameter for each possible classification, C , and combination of our binary variables, 2^D , giving $\mathcal{O}(C2^D)$, the same as the previous case.

Problem C [4 points]: Assume the number of features D is fixed. Let there be N training cases. If the sample size N is very small, which model (Naive Bayes or full) is likely to give lower test set error, and why?

Solution C: For a small N training cases, the Naive Bayes model is likely to generalise better as we avoid the risk of overfitting with fewer parameters. This will likely result in a lower testing error. Alternatively, given we have insufficient data, we are better to assume conditional independence between all features.

Problem D [4 points]: If the sample size N is very large, which model (Naive Bayes or full) is likely to give lower test set error, and why?

Solution D: Inversly, if we have sufficient data, then the risk of overfitting is reduced and we can try to account for the inter-dependence of features for a given input. Trying to limit the number of parameters in the Naive Bayes approach will like result in underfitting, meaning we would expect a lower error in the testing set for the full model.

Problem E [8 EC points]: Assume all the parameter estimates have been computed. What is the computational complexity of making a prediction, i.e. computing $p(y | x)$, using Naive Bayes for a single test case? What is the computation complexity of making a prediction with the full model? In justifying your answer for the full model, choose either the implementation in 1A or 1B and state your choice. For the full-model case, assume that converting a D -bit vector to an array index is an $O(D)$ operation. Also, recall that we have assumed a uniform class prior.

Solution E: We know that:

$$p(y|x) = p(x|y) \frac{p(y)}{p(x)} \quad (4)$$

Where, for the Naive Bayes model, computing $p(x)$ is an $\mathcal{O}(D)$ operation, computing $p(y)$ is an $\mathcal{O}(1)$ operation and computing $p(x|y) = \sum_c p(x|y=c)p(y=c) = \sum_c \prod_d p(x_d|y=c)p(y=c)$ is an $\mathcal{O}(CD)$ operation. Thus, the overall computational time is $\mathcal{O}(CD)$.

For the full model, all the above remains true except for computing $p(x|y)$ where we convert the D -bit vector to an array index ($\mathcal{O}(D)$); after this, it is just a matter of obtaining $p(x|y=c)$ for a corresponding value of y . As such, the overall computational time is $\mathcal{O}(D+C)$.

2 Sequence Prediction [75 Points]

In this problem, we will explore some of the various algorithms associated with Hidden Markov Models (HMMs), as discussed in lecture. We have also uploaded a note on HMMs to the github that might be helpful.

Sequence Prediction

These next few problems will require extensive coding, so be sure to start early!

- You will write an implementation for the hidden Markov model in the cell for HMM Code in the notebook given to you, within the appropriate functions where indicated. There should be no need to write additional functions or use NumPy in your implementation, but feel free to do so if you would like.
- You can (and should!) use the helper cells for each of the subproblems in the notebook, namely 2A, 2Bi, 2Bii, 2C, 2D, and 2F. These can be used to run and check your implementations for each of the corresponding problems. The cells provide useful output in an easy-to-read format. There is no need to modify these cells.
- Lastly, the cell for Utility contains some functions used for loading data directly from the class github repository. There is no need to modify this cell.

The supplementary data folder of the class github repository contains 6 files titled `sequence_data0.txt`, `sequence_data1.txt`, ... , `sequence_data5.txt`. Each file specifies a trained HMM. The first row contains two tab-delimited numbers: the number of states Y and the number of types of observations X (i.e. the observations are $0, 1, \dots, X - 1$). The next Y rows of Y tab-delimited floating-point numbers describe the state transition matrix. Each row represents the current state, each column represents a state to transition to, and each entry represents the probability of that transition occurring. The next Y rows of X tab-delimited floating-point numbers describe the output emission matrix, encoded analogously to the state transition matrix. The file ends with 5 possible emissions from that HMM.

The supplementary data folder also contains one additional file titled `ron.txt`. This is used in problems 2C and 2D and is explained in greater detail there.

Problem A [10 points]: For each of the six trained HMMs, find the max-probability state sequence for each of the five input sequences at the end of the corresponding file. To complete this problem, you will have to implement the Viterbi algorithm (in `viterbi()` of the `HiddenMarkovModel` object). Write your implementation well, as we will be reusing it in a later problem. See the end of problem 2B for a big hint! Note that you do not need to worry about underflow in this part.

In your report, show your results on the 6 files. (Copy-pasting the results of the cell for 2A suffices.)

Solution A: File #0:
Emission Sequence Max Probability State Sequence

25421 31033
01232367534 22222100310
5452674261527433 103100310322222
7226213164512267255 1310331000033100310
0247120602352051010255241 2222222222222222222103

File #1:
Emission Sequence Max Probability State Sequence

77550 22222
7224523677 2222221000
505767442426747 222100003310031
72134131645536112267 10310310000310333100
4733667771450051060253041 222100003222223103222223

File #2:
Emission Sequence Max Probability State Sequence

60622 11111
4687981156 2100202111
815833657775062 0210111111111111
21310222515963505015 02020111111111111021
6503199452571274006320025 1110202111111102021110211

File #3:
Emission Sequence Max Probability State Sequence

13661 00021
2102213421 3131310213
166066262165133 133333133133100
53164662112162634156 20000021313131002133
1523541005123230226306256 1310021333133133133133133

File #4: Emission Sequence Max Probability State Sequence

23664 01124
3630535602 0111201112
350201162150142 011244012441112
00214005402015146362 11201112412444011112
2111266524665143562534450 2012012424124011112411124

File #5:
Emission Sequence Max Probability State Sequence

68535 10111

```
4546566636 1111111111
638436858181213 110111010000011
13240338308444514688 00010000000111111100
0111664434441382533632626 211111111111100111110101
```

Problem B [17 points]: For each of the six trained HMMs, find the probabilities of emitting the five input sequences at the end of the corresponding file. To complete this problem, you will have to implement the Forward algorithm and the Backward algorithm. You may assume that the initial state is randomly selected along a uniform distribution (the starting state transition probabilities are defined in self.A_start in HiddenMarkovModel). Again, write your implementation well, as we will be reusing it in a later problem.

Note that the probability of emitting an input sequence can be found by using either the α vectors from the Forward algorithm or the β vectors from the Backward algorithm. You don't need to worry about this, as it is done for you in probability_alphas() and probability_betas().

Implement the Forward algorithm. In your report, show your results on the 6 files.

Implement the Backward algorithm. In your report, show your results on the 6 files.

After you complete problems 2A and 2B, you can compare your results (the probabilities from the forward and backward algorithms should be the same) for the file titled sequence_data0.txt with the values given in the table below:

Dataset	Emission Sequence	Max-probability State Sequence	Probability of Sequence
0	25421	31033	4.537e-05
0	01232367534	22222100310	1.620e-11
0	5452674261527433	103100310322222	4.348e-15
0	7226213164512267255	1310331000033100310	4.739e-18
0	0247120602352051010255241	222222222222222222103	9.365e-24

Solution B: Forward algorithm:

File #0:

Emission Sequence Probability of Emitting Sequence

```
#####
25421 4.537e-05
01232367534 1.620e-11
5452674261527433 4.348e-15
7226213164512267255 4.739e-18
0247120602352051010255241 9.365e-24
```

File #1:

Emission Sequence Probability of Emitting Sequence

```
#####  
77550 1.181e-04  
7224523677 2.033e-09  
505767442426747 2.477e-13  
72134131645536112267 8.871e-20  
4733667771450051060253041 3.740e-24
```

File #2:

Emission Sequence Probability of Emitting Sequence

```
#####  
60622 2.088e-05  
4687981156 5.181e-11  
815833657775062 3.315e-15  
21310222515963505015 5.126e-20  
6503199452571274006320025 1.297e-25
```

File #3:

Emission Sequence Probability of Emitting Sequence

```
#####  
13661 1.732e-04  
2102213421 8.285e-09  
166066262165133 1.642e-12  
53164662112162634156 1.063e-16  
1523541005123230226306256 4.535e-22
```

File #4:

Emission Sequence Probability of Emitting Sequence

```
#####  
23664 1.141e-04  
3630535602 4.326e-09  
350201162150142 9.793e-14  
00214005402015146362 4.740e-18  
2111266524665143562534450 5.618e-22
```

File #5:

Emission Sequence Probability of Emitting Sequence

```
#####  
68535 1.322e-05  
4546566636 2.867e-09  
638436858181213 4.323e-14  
13240338308444514688 4.629e-18  
0111664434441382533632626 1.440e-22
```

Backward algorithm:

File #0:

Emission Sequence Probability of Emitting Sequence

```
#####
```

25421 4.537e-05
01232367534 1.620e-11
5452674261527433 4.348e-15
7226213164512267255 4.739e-18
0247120602352051010255241 9.365e-24

File #1:

Emission Sequence Probability of Emitting Sequence

77550 1.181e-04
7224523677 2.033e-09
505767442426747 2.477e-13
72134131645536112267 8.871e-20
4733667771450051060253041 3.740e-24

File #2:

Emission Sequence Probability of Emitting Sequence

60622 2.088e-05
4687981156 5.181e-11
815833657775062 3.315e-15
21310222515963505015 5.126e-20
6503199452571274006320025 1.297e-25

File #3:

Emission Sequence Probability of Emitting Sequence

13661 1.732e-04
2102213421 8.285e-09
166066262165133 1.642e-12
53164662112162634156 1.063e-16
1523541005123230226306256 4.535e-22

File #4:

Emission Sequence Probability of Emitting Sequence

23664 1.141e-04
3630535602 4.326e-09
350201162150142 9.793e-14
00214005402015146362 4.740e-18
2111266524665143562534450 5.618e-22

File #5:

Emission Sequence Probability of Emitting Sequence

68535 1.322e-05
4546566636 2.867e-09
638436858181213 4.323e-14


```
13240338308444514688 4.629e-18
0111664434441382533632626 1.440e-22
```

HMM Training

Ron is an avid music listener, and his genre preferences at any given time depend on his mood. Ron's possible moods are happy, mellow, sad, and angry. Ron experiences one mood per day (as humans are known to do) and chooses one of ten genres of music to listen to that day depending on his mood.

Ron's roommate, who is known to take to odd hobbies, is interested in how Ron's mood affects his music selection, and thus collects data on Ron's mood and music selection for six years (2190 data points). This data is contained in the supplementary file `ron.txt`. Each row contains two tab-delimited strings: Ron's mood and Ron's genre preference that day. The data is split into 12 sequences, each corresponding to half a year's worth of observations. The sequences are separated by a row containing only the character `-`.

Problem C [10 points]: Use a single M-step to train a supervised Hidden Markov Model on the data in `ron.txt`. What are the learned state transition and output emission matrices?

Tip: the $(1, 1)$ entry of your transition matrix should be $2.833\text{e-}01$, and the $(1, 1)$ entry of your observation matrix should be $1.486\text{e-}01$.

Solution C: Transition Matrix:

```
#####
2.833e-01 4.714e-01 1.310e-01 1.143e-01
2.321e-01 3.810e-01 2.940e-01 9.284e-02
1.040e-01 9.760e-02 3.696e-01 4.288e-01
1.883e-01 9.903e-02 3.052e-01 4.075e-01
```

Observation Matrix:

```
#####
1.486e-01 2.288e-01 1.533e-01 1.179e-01 4.717e-02 5.189e-02 2.830e-02 1.297e-01 9.198e-02 2.358e-03
1.062e-01 9.653e-03 1.931e-02 3.089e-02 1.699e-01 4.633e-02 1.409e-01 2.394e-01 1.371e-01 1.004e-01
1.194e-01 4.299e-02 6.529e-02 9.076e-02 1.768e-01 2.022e-01 4.618e-02 5.096e-02 7.803e-02 1.274e-01
1.694e-01 3.871e-02 1.468e-01 1.823e-01 4.839e-02 6.290e-02 9.032e-02 2.581e-02 2.161e-01 1.935e-02
```

Problem D [15 points]: Now suppose that Ron has a third roommate who is also interested in how Ron's mood affects his music selection. This roommate is lazier than the other one, so he simply steals the first

roommate's data. Unfortunately, he only manages to grab half the data, namely, Ron's choice of music for each of the 2190 days.

In this problem, we will train an unsupervised Hidden Markov Model on this data. Recall that unsupervised HMM training is done using the Baum-Welch algorithm and will require repeated EM steps. For this problem, we will use 4 hidden states and run the algorithm for 1000 iterations. The transition and observation matrices are initialized for you in the helper functions `supervised_learning()` and `unsupervised_learning()` such that they are random and normalized.

What are the learned state transition and output emission matrices? Please report the result using random seed state 1, as is done by default in the notebook.

Tips for debugging:

- The rows of the state transition and output emitting matrices should sum to 1.
- Your matrices should not change drastically every iteration.
- After many iterations, your matrices should converge.
- If you used random seed 1 for this computation (as is done by default in the notebook), the (1, 1) entry of the state transition matrix should be 5.075e-01, and the (1, 1) entry of the output emission matrix should be 1.117e-01.

Solution D: Transition Matrix:

```
#####
5.075e-01 4.596e-01 6.533e-09 3.292e-02
3.127e-03 2.107e-04 9.964e-01 2.733e-04
1.195e-09 6.886e-02 9.686e-16 9.311e-01
6.203e-01 3.796e-01 1.555e-05 1.579e-04
```

Observation Matrix:

```
#####
1.117e-01 1.525e-01 7.740e-02 1.975e-02 1.594e-01 4.574e-13 3.556e-16 2.475e-01 1.139e-01 1.180e-01
1.205e-01 2.548e-15 1.103e-01 1.751e-01 3.656e-04 2.190e-01 1.002e-01 6.178e-02 1.323e-01 8.053e-02
1.276e-01 2.665e-02 5.788e-02 1.682e-01 1.700e-01 6.969e-02 1.254e-01 3.940e-02 1.627e-01 5.244e-02
1.918e-01 8.206e-02 1.376e-01 8.725e-02 1.152e-01 1.209e-01 1.033e-01 3.101e-02 1.308e-01 5.847e-38
```

Problem E [5 points]: How do the transition and emission matrices from 2C and 2D compare? Which do you think provides a more accurate representation of Ron's moods and how they affect his music choices? Justify your answer. Suggest one way that we may be able to improve the method (supervised or unsupervised) that you believe produces the less accurate representation.

Solution E: The most obvious distinction between the matrices obtained for supervised and unsupervised HMM is that the latter is more-sparse, with values ranging from 10^{-16} to 10^{-1} whilst, in the former, most values are between 10^{-2} and 10^{-1} in the transition matrix. We do note that the supervised model always produces the same transition and observation matrices for a given training set as the minimum is global (obtained from a closed-form solution). On the other hand, by using the Baum-Welch algorithm, unsupervised learning can only find local minima which vary based on our initial starting position. If we assumed the training set adequately represents Ron's mood and the impact on his music choice, then the supervised model is more likely to give an accurate representation since it obtains the global optimum. We have no guarantee that the local minima obtained in the unsupervised model correctly represent his true behaviour (due to hidden states). We could improve our likelihood of ending up in the global minima in our unsupervised model by using more-informed guesses which are closer to the global minima.

Sequence Generation

Hidden Markov Models fall under the umbrella of generative models and therefore can be used to not only predict sequential data, but also to generate it.

Problem F [5 points]: Run the cell for this problem. The code in the cell loads the trained HMMs from the files titled `sequence_data0.txt`, `...`, `sequence_data5.txt` and uses the six models to probabilistically generate five sequences of emissions from each model, each of length 20. In your report, show your results.

Solution F: File #0:

Generated Emission

```
#####  
61372121576400542636  
65257035512152557265  
22503513757650076565  
24302052542075224453  
74164372540666312532
```

File #1:

Generated Emission

```
#####  
67224254254640221044  
53145572605255061365  
03120652632260563554  
57754107154004157145  
50443620741222175553
```

File #2:

Generated Emission

```
#####  
23492277012312492009  
92715253263463500777  
00206169901079655318  
34658526761560627594  
98316963019608446826
```

File #3:

Generated Emission

```
#####  
53324006315616155155  
61131140351050201522  
26364625661106214604  
02325053164160635305  
20066023620261022165
```

File #4:

Generated Emission

```
#####  
63333124051064432222  
54046611655614306333  
30164016126063513444  
66156363216364120060  
44664503206213132626
```

File #5:

Generated Emission

```
#####  
23118681064868303466  
68845455633644828521  
30418103230140186766  
10567813866403462083  
82316383453840668047
```

Visualization & Analysis

Once you have implemented the HMM code part of the notebook, load and run the cells for the following subproblems. Here you will apply the HMM you have implemented to the Constitution. There is no coding required for this part, only analysis.

Answer the following problems in the context of the visualizations in the notebook.

Problem G [3 points]: What can you say about the sparsity of the trained A and O matrices? How does this sparsity affect the transition and observation behaviour at each state?

Solution G: Although both are sparse, O is generally more-sparse than A . A zero element in row i of matrix A indicates that the state i has a zero probability of transitioning to the state corresponding to that column (thus, a sparse row i means that there are fewer states to transition to). Similarly, a zero element in row i of matrix O indicates a zero probability of generating the observation corresponding to that column from state i ; thus a sparse row i indicates that few observations can be generated from state i . If all elements in a column of O are very small valued, it indicates that this observation is very unlikely to be generated. As O is more-sparse than A , it indicates that the relationship between states and observations is stronger than it is between states (transitioning).

Problem H [5 points]: How do the sample emission sentences from the HMM change as the number of hidden states is increased? What happens in the special case where there is only one hidden state? In general, when the number of hidden states is unknown while training an HMM for a fixed observation set, can we increase the training data likelihood by allowing more hidden states?

Solution H: As we increase the number of hidden states, the sample emission begin to resemble something akin to an actual sentence from the constitution. In the case when there is only one hidden layer, the sample emission sentence doesn't make any sense; this is because the words are picked completely at random given, with one hidden layer, all words have the same probability of being picked (weighted by the frequency at which they occur in the original text). This partly explains why the word 'president' appears twice.

We can indeed increase training data likelihood by increasing the number of allowed hidden states. Doing so results in sparser transition and observation matrices meaning there is a stronger relationship between states and observations, and other states. In the extreme case where we have the same number of states as observations, our observation matrix should approach some sort of identity matrix (with some columns permuted), meaning each state is perfectly associated with one observation, maximising our training likelihood.

Problem I [5 points]: Pick a state that you find semantically meaningful, and analyze this state and its wordcloud. What does this state represent? How does this state differ from the other states? Back up your claim with a few key words from the wordcloud.

Solution I: State 3 appears to be the most semantically meaningful, where most of the words in this state are either numbers (two, three, four, seven, eight, thirty, etc.) and words that usually involve counting (days, ballot, grant, pay, etc.). Most of the other states appear to be semi-random, linked to the frequency of the word or words just linked to the constitution. Most of the words in state 3 serve very similar grammatical purposes.